

# Aim: To implement a Machine Learning Classification model using a Logistic regression algorithm.

```
In [11]: import pandas as pd
from matplotlib import pyplot as plt
%matplotlib inline
import seaborn as sns
import plotly.express as px
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.metrics import precision_recall_fscore_support
from sklearn import metrics
```

```
In [12]: df = pd.read_csv(r"C:\Users\VICTUS\Desktop\Sem 6 Practicals\ML\Diabetes.csv")
df.head()
```

```
Out[12]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeF
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

```
In [13]: df.isna().sum()
```

```
Out[13]: Pregnancies      0
Glucose      0
BloodPressure  0
SkinThickness  0
Insulin      0
BMI          0
DiabetesPedigreeFunction  0
Age          0
Outcome      0
dtype: int64
```

```
In [14]: df.info
```

```
Out[14]: <bound method DataFrame.info of
hickness  Insulin  BMI  \
0          6    148    72    35    0  33.6
1          1     85    66    29    0  26.6
2          8    183    64     0    0  23.3
3          1     89    66    23   94  28.1
4          0    137    40    35  168  43.1
..      ...    ...    ...    ...    ...
763        10    101    76    48  180  32.9
764         2    122    70    27    0  36.8
765         5    121    72    23  112  26.2
766         1    126    60     0    0  30.1
767         1     93    70    31    0  30.4

DiabetesPedigreeFunction  Age  Outcome
0          0.627    50         1
1          0.351    31         0
2          0.672    32         1
3          0.167    21         0
4          2.288    33         1
..      ...    ...    ...
763        0.171    63         0
764        0.340    27         0
765        0.245    30         0
766        0.349    47         1
767        0.315    23         0

[768 rows x 9 columns]>
```

```
In [15]: df.describe()
```

Out[15]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000

```
In [16]: x = df.iloc[:,[0,7]].values
y = df.iloc[:, -1].values
```

```
In [17]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state
```

```
In [19]: classifier = LogisticRegression()
classifier.fit(x_train,y_train)
```

```
Out[19]: ▾ LogisticRegression
LogisticRegression()
```

```
In [20]: y_pred= classifier.predict(x_test)
```

```
In [21]: print(y_test)
```

```
[1 0 0 1 0 0 1 1 0 0 1 1 0 0 0 0 1 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1
 0 0 0 0 0 0 1 1 0 0 1 1 1 0 0 1 0 0 0 0 1 1 1 1 0 0 1 1 1 0 0 0 0 0 0 0
 1 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 1 0 1 1 0 0 0 0 0 1 0 0 0 1 0
 1 1 1 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 1 0 1 1 0 0 0 0 0 1 0 0 0
 0 1 0 1 0 0 1 0 0 0 1 1 1 1 0 0 0 1 0 0 0 0 0 0 1 1 0 0 0 0 0 0 1 1 0 1 1
 0 1 1 1 0 0 0]
```

```
In [22]: print(y_pred)
```

```
[0 0 0 1 0 0 0 0 0 1 0 1 1 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 1
 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0
 0 0 0 0 1 0 0 0 1 0 1 0 1 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 1 0 0 0 0 0 1 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 1 0 0 0 0 1 0 0
 0 0 0 0 0 1 0]
```

## Confusion Matrix

```
In [24]: cm = confusion_matrix(y_test,y_pred)
print(cm)
```

```
[[117  13]
 [ 49  13]]
```

```
In [25]: precision_recall_fscore_support(y_test,y_pred, average='macro')
```

```
Out[25]: (0.6024096385542168, 0.5548387096774193, 0.542997542997543, None)
```

```
In [26]: precision_recall_fscore_support(y_test,y_pred, average='micro')
```

```
Out[26]: (0.6770833333333334, 0.6770833333333334, 0.6770833333333334, None)
```

```
In [27]: precision_recall_fscore_support(y_test,y_pred, average='weighted')
```

```
Out[27]: (0.638679718875502, 0.6770833333333334, 0.6306690212940212, None)
```

```
In [28]: accuracy_score(y_test,y_pred)
```

```
Out[28]: 0.6770833333333334
```

```
In [29]: Accuracy = metrics.accuracy_score(y_test,y_pred)
Accuracy
```


```
Out[29]: 0.6770833333333334
```

```
In [30]: classifier.intercept_
classifier.coef_
```

```
Out[30]: array([[0.05180623, 0.03442028]])
```

```
In [32]: plt.figure(figsize=(9,9))  
sns.heatmap(cm, annot=True, fmt='d', cbar=False)  
plt.xlabel('Predicted Labels')  
plt.ylabel('True labels')  
plt.show
```

```
Out[32]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [36]: precision, recall, f1_score, _ = precision_recall_fscore_support(y_test, y_pred)  
print('Precision: ', precision)  
print('Recall: ', recall)  
print('F1_score: ', f1_score)
```

```
Precision: [0.70481928 0.5      ]  
Recall:    [0.9      0.20967742]  
F1_score:  [0.79054054 0.29545455]
```